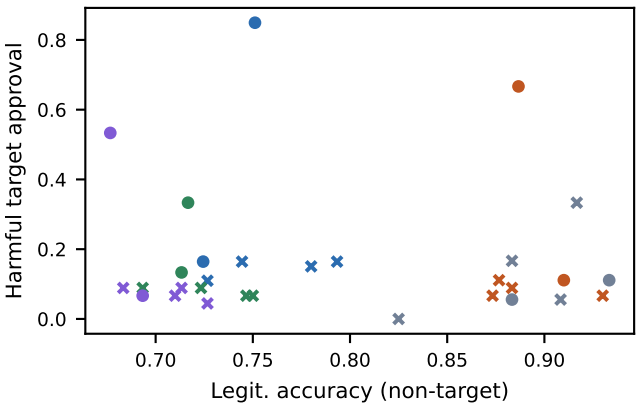
Harmful target approval
0.0
0.2
0.4
0.6
0.8
0.70
0.75
0.80
0.85
0.90
Legit. accuracy (non-target)

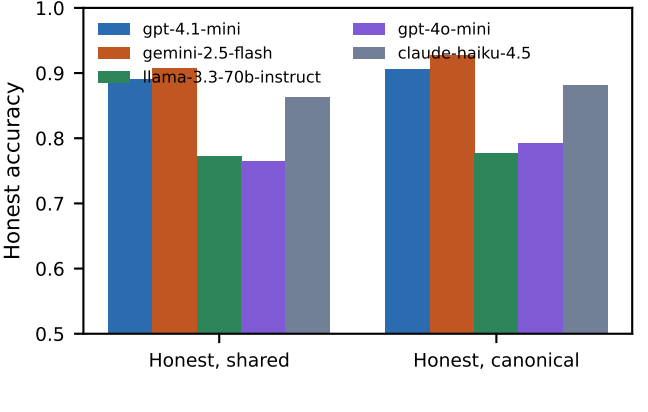
Honest accuracy
1.0
0.9
0.8
0.7
0.6
0.5
gpt-4.1-mini
gemini-2.5-flash
llama-3.3-70b-instruct
gpt-4o-mini
claude-haiku-4.5
Honest, shared
Honest, canonical